

Data Mining Note

Chauncey Liu

November 13, 2017

Contents

1	Association Rules and Frequent Pattern Mining	2
1.1	The Frequent Pattern Mining Model	2
1.2	Association Rules	3
1.2.1	Mining Association Rules	3
1.2.2	Itemset Lattice	3
1.2.3	Anti-monotonicity (downward closure) property	3
1.2.4	Computing the Association Rules	3
1.2.5	Interestingness of Association Rules	4
1.3	The Apriori Algorithm	4
1.3.1	Approach	4
1.3.2	Candidate Generation	5
1.3.3	Support Counting	5
1.3.4	Hash Tree	5
1.3.5	Methods of Efficiency Improvement	6
1.4	Enumeration-Tree Algorithms	6
1.5	Suffix-based Pattern Growth Methods	6
1.6	Constraint-Based Association Mining	6
1.7	Multi-level Association Rules	6
1.8	Pattern Summarization	6

1 Association Rules and Frequent Pattern Mining

1.1 The Frequent Pattern Mining Model

Basic Concepts:

- Items $I = \{i_1, i_2, \dots, i_m\}$, a set of literals
- Itemset X (transaction): set of items $X \subseteq I$
- Database D : set of transactions T_i where $T_i \subseteq I$
- T contains X : $X \subseteq T$
- Items in transactions or itemsets are sorted in lexicographic order
- Length of itemset: number of elements of itemset
- k-itemset: itemset of length k

Supermarket example:

- Items $I = \{Bread, Butter, Milk, Eggs, Yogurt\}$
- Itemset X , Transaction T :
 - $X_1 = \{Bread, Butter\}$, $X_2 = \{Eggs, Milk\}$
 - $T_1 = \{Bread, Butter, Milk\}$, $T_2 = \{Eggs, Milk, Yogurt\}$
- Database D : set of transactions T_i where $T_i \subseteq I$
 - $T_1 = \{Bread, Butter, Milk\}$, $T_2 = \{Eggs, Milk, Yogurt\}$
 - $D = \{T_1, T_2\} = \{\{Bread, Butter, Milk\}, \{Eggs, Milk, Yogurt\}\}$
- T contains X : $X \subseteq T$
 - $X_1 = \{Bread, Butter\}$, $T_1 = \{Bread, Butter, Milk\}$
 - $X_1 \subseteq T_1$
- Items in transactions or itemsets are sorted in lexicographic order
- Length of itemset: number of elements of itemset
- k-itemset: itemset of length k

Further Concept

- Support of itemset X in D : percentage of transactions in D containing X
 - $sup(X, D) = \frac{|T \in D | X \subseteq T|}{|D|}$
- Frequent itemset X in D : item set X with
 - $freq(X, D) :\Leftrightarrow sup(X, D) \geq minsup$
- Association rule: implication of the form $X \Rightarrow Y$
 - where $X \subseteq I, Y \subseteq I$ and $X \cap Y = \emptyset$
- Support s of association rule $X \Rightarrow Y$ in D :
 - indicates how frequently the itemset appears in the dataset
 - support of $X \cup Y$ in D
 - $s = \frac{|T \in D | (X \cup Y) \subseteq T|}{|D|}$
- Confidence c of association rule $X \Rightarrow Y$ in D :
 - indicates how often the rule has been found to be true
 - percentage of transactions containing Y in the subset of all transactions in D that contain X
 - $c = \frac{|T \in D | (X \cup Y) \subseteq T|}{|T \in D | X \subseteq T|} = \frac{sup(X \cup Y)}{sup(X)}$

1.2 Association Rules

1.2.1 Mining Association Rules

1. Determine the frequent itemsets in the database

Naive algorithm: count the frequencies of all k -itemsets $\subseteq I$, inefficient since $\binom{m}{k}$ such itemsets

2. Generate the association rules from the frequent itemsets

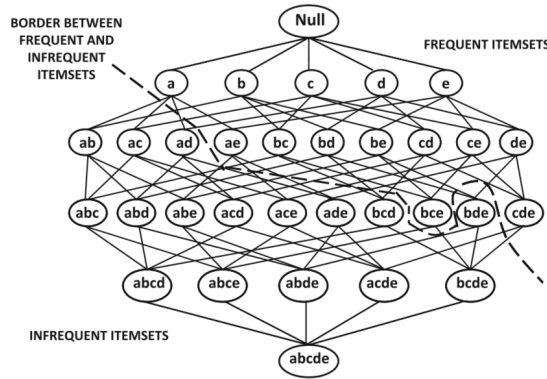
Itemset X frequent and $A \subseteq X$

$A \Rightarrow (X - A)$ satisfies minimum support constraint (confidence remaining to be checked)

1.2.2 Itemset Lattice

- Lattice: a partially ordered set with unique least upper bound and greatest lower bound.
- Itemset lattice:
 - elements: itemsets $X_1 \subseteq I, X_2 \subseteq I, \dots, X_n \subseteq I$
 - partial order: $X_1 < X_2 :\Leftrightarrow X_1 \subset X_2$
 - least upper bound: I
 - greatest lower bound: \emptyset

Figure 1: Itemset Lattice



1.2.3 Anti-monotonicity (downward closure) property

Each subset of a frequent itemset is also frequent.

$$\forall T_1 \subseteq I, T_2 \subseteq I : T_1 \subseteq T_2 \wedge freq(T_2, D) \Rightarrow freq(T_1, D)$$

because of

$$\forall T_1 \subseteq I, T_2 \subseteq I : T_1 \subseteq T_2 \Rightarrow sup(T_1, D) \geq sup(T_2, D)$$

If one subset is not frequent, then superset cannot be frequent.

This property makes frequent itemset mining efficient, since in practice most itemsets are infrequent.

1.2.4 Computing the Association Rules

- Given a frequent itemset X
- For each subset A of X , form the rule $A \Rightarrow (X - A)$
- Compute confidence of the rule $A \Rightarrow (X - A)$
 - $confidence(A \Rightarrow (X - A)) = \frac{sup(X)}{sup(A)}$
- Discard rules that do not have minimum confidence
- Store frequent itemsets with their supports in a hash table
 - no DB accesses, no disk I/O

1.2.5 Interestingness of Association Rules

- Filter out misleading association rules
- Expected support for the rule $A \Rightarrow B$
 - $P(A \cup B) = P(A) \cdot P(B)$
 - assuming the independence of A and B
- Interestingness measure for rule $A \Rightarrow B$
 - $\frac{P(A \cup B)}{P(A)} - P(B)$
 - The larger this measure, the more interesting the discovered association between A and B
- An alternative interestingness measure is the lift of an association rule
- If A and B are independent, then $P(A \cup B) = P(A) \cdot P(B)$
 - i.e. $\frac{P(A \cup B)}{P(A) \cdot P(B)} = 1$
- We define the lift of a rule $A \Rightarrow B$ as follows:
 - $lift(A \Rightarrow B) = \frac{P(A \cup B)}{P(A) \cdot P(B)}$
 - Can also be formulated as:
 - * $lift(A \Rightarrow B) = \frac{P(A \cup B)/P(A)}{P(B)} = \frac{support_{actual}}{support_{expected}}$
 - * as the ratio of the conditional probability $P(B|A)$ and the unconditional probability $P(B)$
 - A lift $\gg 1$ indicates that the discovered association between A and B is interesting.

1.3 The Apriori Algorithm

1.3.1 Approach

- Determine first the frequent 1-itemsets, then frequent 2-itemsets, ...
- To determine the frequent k+1-items, consider only the k+1-items for which all k-subsets are frequent
- Calculation of support: one database scan counting the support for all relevant itemsets

Algorithm 1 Algorithm Apriori

/* C_k : set of candidate itemsets of length k */
 /* F_k : set of all frequent itemsets of length k */

```

function APRIORI(D, minsup)
  while  $F_k \neq \emptyset$  do
    Generate  $C_{k+1}$  by joining itemset-pairs in  $F_k$ ;
    Prune itemsets from  $C_{k+1}$  that violate anti-monotonicity;
    Determine  $F_{k+1}$  by counting support of  $C_{k+1}$  in D and retaining itemsets from  $C_{k+1}$  with support
    at least minsup;
     $k = k + 1$ ;
  return  $\cup_k F_k$ ;
  
```

1.3.2 Candidate Generation

Requirements for set C_k of candidate itemsets

- Superset of F_k
- Significantly smaller than set of all k -subsets of I

Step1: Join

- Frequent $k-1$ -itemsets p and q , p and q are joined if they agree in their first $k-2$ items
- E.g. $p \in F_{k-1} = \{1, 2, 3\}$, $q \in F_{k-1} = \{1, 2, 4\} \Rightarrow (1, 2, 3, 4) \in C_k$
- Choose first $k-2$ items to avoid duplication without missing any candidates

Step2: Pruning

- Remove all elements from C_k having a $k-1$ -subset not contained in F_{k-1}
- E.g. $F_3 = \{(1, 2, 3), (1, 2, 4), (1, 3, 4), (1, 3, 5), (2, 3, 4)\}$
- After join step: $C_4 = \{(1, 2, 3, 4), (1, 3, 4, 5)\}$
- In pruning step: remove $(1, 3, 4, 5)$ since subsets $(1, 4, 5), (3, 4, 5)$ are missing
- $C_4 = \{(1, 2, 3, 4)\}$

1.3.3 Support Counting

for each candidate $c \in C_k$ **do** $c.count = 0$;

for each transaction $T \in D$ **do**

$CT := subset(C_k, T)$; // all candidates from C_k that are contained in transaction T

for each candidate $c \in CT$ **do** $c.count++$;

$F_k := \{c \in C_k | (c.count/|D|) \geq minsup\}$

To achieve one scan over the database D , $subset(C_k, T)$ should be implemented properly. Thus we need Hash Tree.

1.3.4 Hash Tree

Hash tree as a data structure for C_k

- Leaf node: records list of itemsets (with frequencies)
- Inner node: contains hash table (apply hash function to d -th elements), each hash bucket at level d references son node at level $d+1$
- Root has level 1

Finding an itemset

- Start from the root
- At level d : apply hash function h to the d -th element of the itemset.

Inserting an itemset

- Find the corresponding leaf node and insert new itemset
- In case of overflow:
 - Covert leaf node into inner node and create all its son nodes (new leaves).
 - Distribute all entries over the new leaf nodes according to hash function h .

Find all candidates contained in $T = (t_1 t_2 t_3 \dots t_m)$

- At root
 - Determine hash values $h(t_i)$ for each item t_i in T
 - Continue search in all corresponding son nodes
- At inner node of level d
 - Assumption: inner node has been reached by hashing t_i
 - Determine hash values and continue search for all items t_k in T with $k > i$
- At leaf node
 - For each itemset X in this node, test whether $X \subseteq T$

1.3.5 Methods of Efficiency Improvement

- Support counting using a hash table
 - Hash table instead of hash tree, support counters for hash buckets
 - k-itemset with corresponding bucket counter \downarrow minsup cannot be frequent
 - * more efficient access to candidates but inaccurate counts
- Reduction of transactions
 - Transactions that do not contain any frequent k-itemset are irrelevant
 - Remove such transactions for future phases
 - * more efficient database scan, but additional writing of database
- Partitioning of the database
 - Itemset is only frequent if frequent in at least one partition
 - Form memory-resident partitions of the database
 - * more efficient on partitions, but expensive combination of intermediate results
- Sampling
 - Apply algorithm to sample to find frequent itemsets
 - Count support of these frequent itemsets in the whole database
 - Determine further candidates and support counting on the whole database

1.4 Enumeration-Tree Algorithms

1.5 Suffix-based Pattern Growth Methods

1.6 Constraint-Based Association Mining

1.7 Multi-level Association Rules

1.8 Pattern Summarization